

▼ Import Libraries

```
1 import pandas as pd
2 import seaborn as sns
3 import matplotlib.pyplot as plt
4 import numpy as np
5 sns.set_theme(color_codes=True)
```

▼ Classification Supervised Learning

You can get the dataset from here : [Diabetes Dataset](#)

```
1 df = pd.read_csv('diabetes.csv')
2 df.head()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigree
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	

Split the X and y then split the train and test data into like this :

Pregnancies - Age	Outcome
X train	y train
X test	y test

```
1 X = df.drop('Outcome', axis=1)
2 y = df['Outcome']
```

```
1 #test size 10% and train size 90%
2 #Random state = Controls the randomness of the bootstrapping of the samples
3 from sklearn.model_selection import train_test_split
4 from sklearn.metrics import accuracy_score
5 X_train, X_test, y_train, y_test = train_test_split(X,y, test_size=0.1, random_state=0)
```

▼ Decision Tree Classifier

```
1 from sklearn.tree import DecisionTreeClassifier
2
3 dtree = DecisionTreeClassifier(random_state=0) #Controls the randomness of the bootstrapping of the samples
4 dtree.fit(X_train, y_train)
```

▼ DecisionTreeClassifier
DecisionTreeClassifier(random_state=0)

```
1 y_pred = dtree.predict(X_test)
2 print("Accuracy Score :", round(accuracy_score(y_test, y_pred)*100 ,2), "%")
```

Accuracy Score : 75.32 %

```
1 from sklearn.metrics import accuracy_score, f1_score, precision_score, recall_score
2 print('F-1 Score : ',(f1_score(y_test, y_pred)))
3 print('Precision Score : ',(precision_score(y_test, y_pred)))
4 print('Recall Score : ',(recall_score(y_test, y_pred)))
```

F-1 Score : 0.6666666666666667
Precision Score : 0.6129032258064516
Recall Score : 0.7307692307692307

▼ Random Forest Classifier

```
1 from sklearn.ensemble import RandomForestClassifier
2 rfc = RandomForestClassifier(random_state=0)
3 rfc.fit(X_train, y_train)
```

▼ RandomForestClassifier
RandomForestClassifier(random_state=0)

```
1 y_pred = rfc.predict(X_test)
2 print("Accuracy Score :", round(accuracy_score(y_test, y_pred)*100 ,2), "%")
```

Accuracy Score : 80.52 %

```
1 from sklearn.metrics import accuracy_score, f1_score, precision_score, recall_score
2 print('F-1 Score : ',(f1_score(y_test, y_pred)))
3 print('Precision Score : ',(precision_score(y_test, y_pred)))
4 print('Recall Score : ',(recall_score(y_test, y_pred)))
```

F-1 Score : 0.7058823529411765
Precision Score : 0.72
Recall Score : 0.6923076923076923

▼ Support Vector Machine Classifier

```
1 from sklearn import svm
2 svm = svm.SVC(random_state=0)
3 svm.fit(X_train, y_train)
```

▼ SVC
SVC(random_state=0)

```
1 y_pred = svm.predict(X_test)
2 print("Accuracy Score :", round(accuracy_score(y_test, y_pred)*100 ,2), "%")
```

Accuracy Score : 85.71 %

```
1 from sklearn.metrics import accuracy_score, f1_score, precision_score, recall_score
2 print('F-1 Score : ',(f1_score(y_test, y_pred)))
3 print('Precision Score : ',(precision_score(y_test, y_pred)))
4 print('Recall Score : ',(recall_score(y_test, y_pred)))
```

F-1 Score : 0.7659574468085107
Precision Score : 0.8571428571428571

Recall Score : 0.6923076923076923

▼ AdaBoost Classifier

```
1 from sklearn.ensemble import AdaBoostClassifier
2 ada = AdaBoostClassifier(random_state=0)
3 ada.fit(X_train, y_train)
```

▼ AdaBoostClassifier
AdaBoostClassifier(random_state=0)

```
1 y_pred = ada.predict(X_test)
2 print("Accuracy Score :", round(accuracy_score(y_test, y_pred)*100 ,2), "%")
```

Accuracy Score : 79.22 %

```
1 from sklearn.metrics import accuracy_score, f1_score, precision_score, recall_score
2 print('F-1 Score : ',(f1_score(y_test, y_pred)))
3 print('Precision Score : ',(precision_score(y_test, y_pred)))
4 print('Recall Score : ',(recall_score(y_test, y_pred)))
```

F-1 Score : 0.7037037037037038
Precision Score : 0.6785714285714286
Recall Score : 0.7307692307692307

▼ Logistic Regression Classifier

```
1 from sklearn.linear_model import LogisticRegression
2 lr = LogisticRegression()
3 lr.fit(X_train, y_train)
```

/usr/local/lib/python3.9/dist-packages/sklearn/linear_model/_logistic.py:458: Conv
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

n_iter_i = _check_optimize_result(

▼ LogisticRegression
LogisticRegression()



```
1 y_pred = lr.predict(X_test)
2 print("Accuracy Score :", round(accuracy_score(y_test, y_pred)*100 ,2), "%")
```

Accuracy Score : 87.01 %

```
1 from sklearn.metrics import accuracy_score, f1_score, precision_score, recall_score
2 print('F-1 Score : ',(f1_score(y_test, y_pred)))
3 print('Precision Score : ',(precision_score(y_test, y_pred)))
4 print('Recall Score : ',(recall_score(y_test, y_pred)))
```

F-1 Score : 0.7916666666666666
Precision Score : 0.8636363636363636
Recall Score : 0.7307692307692307

▼ Regression Supervised Learning

You can get the dataset from here : [Mobile Price Dataset](#)

```
1 df2 = pd.read_csv('Cellphone.csv')
2 df2.head()
```

	Product_id	Price	Sale	weight	resolution	ppi	cpu core	cpu freq	internal mem	ram
0	203	2357	10	135.0	5.2	424	8	1.35	16.0	3.000
1	880	1749	10	125.0	4.0	233	2	1.30	4.0	1.000
2	40	1916	10	110.0	4.7	312	4	1.20	8.0	1.500
3	99	1315	11	118.5	4.0	233	2	1.30	4.0	0.512
4	880	1749	11	125.0	4.0	233	2	1.30	4.0	1.000

```

1 #Remove Product_id attribute because it's unnecessary
2 df2.drop(columns='Product_id', inplace = True)
3 df2.head()

```

	Price	Sale	weight	resolution	ppi	cpu core	cpu freq	internal mem	ram	RearCam	Fro
0	2357	10	135.0	5.2	424	8	1.35	16.0	3.000	13.00	
1	1749	10	125.0	4.0	233	2	1.30	4.0	1.000	3.15	
2	1916	10	110.0	4.7	312	4	1.20	8.0	1.500	13.00	
3	1315	11	118.5	4.0	233	2	1.30	4.0	0.512	3.15	

Just like the Classification method, split the target and attribute then split the train and test data

```

1 X = df2.drop('Price', axis=1)
2 y = df2['Price']

```

```

1 #test size 20% and train size 80%
2 from sklearn.model_selection import train_test_split
3 from sklearn.metrics import accuracy_score
4 X_train, X_test, y_train, y_test = train_test_split(X,y, test_size=0.2,random_state=0)

```

Decision Tree Regressor

```

1 from sklearn.tree import DecisionTreeRegressor
2 dtree = DecisionTreeRegressor(random_state=0)
3 dtree.fit(X_train, y_train)

```

DecisionTreeRegressor
DecisionTreeRegressor(random_state=0)

```

1 from sklearn import metrics
2 import math
3 y_pred = dtree.predict(X_test)
4 mae = metrics.mean_absolute_error(y_test, y_pred)
5 mse = metrics.mean_squared_error(y_test, y_pred)
6 r2 = metrics.r2_score(y_test, y_pred)
7 rmse = math.sqrt(mse)
8
9 print('MAE is {}'.format(mae))
10 print('MSE is {}'.format(mse))
11 print('R2 score is {}'.format(r2))
12 print('RMSE score is {}'.format(rmse))

```

```

MAE is 53.63636363636363
MSE is 19068.666666666668
R2 score is 0.9593394123291379
RMSE score is 138.0893430597259

```

▼ Random Forest Regressor

```
1 from sklearn.ensemble import RandomForestRegressor
2 rf = RandomForestRegressor(random_state=0)
3 rf.fit(X_train, y_train)
```

```
▼ RandomForestRegressor
RandomForestRegressor(random_state=0)
```

```
1 from sklearn import metrics
2 import math
3 y_pred = rf.predict(X_test)
4 mae = metrics.mean_absolute_error(y_test, y_pred)
5 mse = metrics.mean_squared_error(y_test, y_pred)
6 r2 = metrics.r2_score(y_test, y_pred)
7 rmse = math.sqrt(mse)
8
9 print('MAE is {}'.format(mae))
10 print('MSE is {}'.format(mse))
11 print('R2 score is {}'.format(r2))
12 print('RMSE score is {}'.format(rmse))
```

```
MAE is 107.12181818181816
MSE is 17579.776721212114
R2 score is 0.962514209037145
RMSE score is 132.58875035692927
```

▼ AdaBoost Regressor

```
1 from sklearn.ensemble import AdaBoostRegressor
2 ada = AdaBoostRegressor(random_state=0)
3 ada.fit(X_train, y_train)
```

```
▼ AdaBoostRegressor
AdaBoostRegressor(random_state=0)
```

```
1 from sklearn import metrics
2 import math
3 y_pred = ada.predict(X_test)
4 mae = metrics.mean_absolute_error(y_test, y_pred)
5 mse = metrics.mean_squared_error(y_test, y_pred)
6 r2 = metrics.r2_score(y_test, y_pred)
7 rmse = math.sqrt(mse)
8
9 print('MAE is {}'.format(mae))
10 print('MSE is {}'.format(mse))
11 print('R2 score is {}'.format(r2))
12 print('RMSE score is {}'.format(rmse))
```

```
MAE is 130.3148744244737
MSE is 24064.750340734758
R2 score is 0.9486861400373987
RMSE score is 155.1281739102693
```

✓ 0s completed at 1:57 AM

● ×