

5 Flower Types Classification Dataset

The **5 Flower Types Classification Dataset** is a collection of images belonging to five different flower classes: Lilly, Lotus, Sunflower, Orchid, and Tulip. Each flower class contains 1000 images, resulting in a total of 5000 images in the dataset.

This dataset is suitable for training and evaluating a multi-class Convolutional Neural Network (CNN) model to classify flower images into one of the five mentioned classes. The goal of the classification task is to accurately identify the type of flower from an input image.

The dataset can be used to explore various deep learning techniques for image classification, such as data augmentation, transfer learning, and model fine-tuning. It provides a challenging task due to the visual similarity and subtle differences among different flower types.

Dataset Details:

- Number of classes: 5
- Total images: 5000 (1000 images per class)
- Image format: JPG or PNG
- Image resolution: Varies (please preprocess the images to a consistent size if required)

The 5 Flower Types Classification Dataset is a valuable resource for researchers, students, and practitioners interested in the field of computer vision, specifically in image classification tasks. It can be used for educational purposes, benchmarking different models, and advancing the state-of-the-art in flower classification.

Feel free to download the dataset and start exploring the fascinating world of flower image classification!:

<https://www.kaggle.com/datasets/kausthubkannan/5-flower-types-classification-dataset>

<https://www.kaggle.com/datasets/kausthubkannan/5-flower-types-classification-dataset>

```
In [ ]: 1 # import libraries
        2 import os
        3 import shutil
        4 import random
        5 import numpy as np
        6 import pandas as pd
        7 from matplotlib import pyplot as plt
        8 from PIL import Image
        9 from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
       10 import tensorflow as tf
       11 from tensorflow.keras.preprocessing.image import ImageDataGenerator
       12 from tensorflow.keras.optimizers import Adam
       13 from tensorflow.keras.layers import Conv2D, MaxPool2D, Flatten, Dense, Dropout
       14 from tensorflow.keras import Sequential
       15 os.environ['KAGGLE_CONFIG_DIR'] = '/content'
```

Downloading and preparing data for model

```
In [ ]: 1 !kaggle datasets download -d kausthubkannan/5-flower-types-classification-dataset
```

Warning: Your Kaggle API key is readable by other users on this system! To fix this, you can run 'chmod 600 /content/kaggle.json'

Downloading 5-flower-types-classification-dataset.zip to /content

96% 232M/242M [00:02<00:00, 104MB/s]

100% 242M/242M [00:02<00:00, 115MB/s]

```
In [ ]: 1 !unzip \*.zip && rm *.zip
```

Streaming output truncated to the last 5000 lines.

```
inflating: flower_images/Lilly/00048a5c76.jpg
inflating: flower_images/Lilly/001ff6644e.jpg
inflating: flower_images/Lilly/001ff6656j.jpg
inflating: flower_images/Lilly/00973ad1b1.jpg
inflating: flower_images/Lilly/00a7d512d6.jpg
inflating: flower_images/Lilly/00f36a3c40.jpg
inflating: flower_images/Lilly/013628cccc.jpg
inflating: flower_images/Lilly/01998d6fb5.jpg
inflating: flower_images/Lilly/01a0ec319c.jpg
inflating: flower_images/Lilly/01b4bb0289.jpg
inflating: flower_images/Lilly/025ef3ea44.jpg
inflating: flower_images/Lilly/02a7a2df46.jpg
inflating: flower_images/Lilly/02be2ca388.jpg
inflating: flower_images/Lilly/035cce082f.jpg
inflating: flower_images/Lilly/039eba79d4.jpg
inflating: flower_images/Lilly/04067b91d6.jpg
inflating: flower_images/Lilly/04acfd5449.jpg
inflating: flower_images/Lilly/05777790e2.jpg
inflating: flower_images/Lilly/05f344160c.jpg
```

```
In [ ]: 1 !pip install split-folders
```

Looking in indexes: <https://pypi.org/simple>, (<https://pypi.org/simple>,) <https://us-python.pkg.dev/colab-wheels/public/simple/> (<https://us-python.pkg.dev/colab-wheels/public/simple/>)
Collecting split-folders
 Downloading split_folders-0.5.1-py3-none-any.whl (8.4 kB)
Installing collected packages: split-folders
Successfully installed split-folders-0.5.1

```
In [ ]: 1 import splitfolders
```

```
In [ ]: 1 src_dir = '/content/flower_images'
2 dst_dir = '/content/Data'
```

Data preprocessing (image augmentation)

```
In [ ]: 1 splitfolders.ratio(input=src_dir, output=dst_dir, ratio=(0.8, 0.2))
```







Copying files: 5000 files [00:00, 5625.05 files/s]

```
In [ ]: 1 train_datagen = ImageDataGenerator(rescale=1/255., rotation_range=0.2,
2                                           # brightness_range=(0.2, 0.5),
3                                           zoom_range=0.2, shear_range=0.2,
4                                           horizontal_flip=True)
5 train_dataset = train_datagen.flow_from_directory('/content/Data/train',
6                                                    target_size=(150, 150),
7                                                    batch_size=32,
8                                                    shuffle=True)
9
10
11 val_datagen = ImageDataGenerator(rescale=1/255.)
12 val_dataset = val_datagen.flow_from_directory('/content/Data/val', target_size=(150, 150),
13                                              batch_size=32, shuffle=False)
```

Found 4000 images belonging to 5 classes.
Found 1000 images belonging to 5 classes.

In []:

```
1 images, labels = next(train_dataset)
2 labels = np.argmax(labels, axis=1)
3 class_names = list(train_dataset.class_indices.keys())
4 def plot_random_images(images, labels, class_names):
5     plt.figure(figsize=(12, 6))
6
7     for i in range(6):
8         ax = plt.subplot(2, 3, i+1)
9         rand_index = random.choice(range(len(images)))
10        plt.imshow(images[rand_index])
11        plt.title(class_names[labels[rand_index]], color='green', fontsize=12)
12
13    plt.tight_layout()
14    plt.show()
15
16 plot_random_images(images, labels, class_names)
```



creating model

```
In [ ]: 1 model = Sequential([
2
3         Conv2D(filters=16, kernel_size=(3,3), strides=1, activation='relu', input_s
4         MaxPool2D(pool_size=(2,2), strides=2, padding='valid'),
5
6         Conv2D(filters=32, kernel_size=(3,3), strides=2, activation='relu'),
7         MaxPool2D(pool_size=(2,2), strides=1, padding='same'),
8
9         Conv2D(filters=64, kernel_size=(3,3), strides=2, activation='relu'),
10        MaxPool2D(pool_size=(2,2), strides=1, padding='same'),
11
12        Flatten(),
13        Dense(256, activation='relu'),
14        Dense(128, activation='relu'),
15        Dense(64, activation='relu'),
16        Dense(5, activation='softmax')
17
18    ])
```

```
In [ ]: 1 model.compile(optimizer=Adam(), loss=tf.keras.losses.CategoricalCrossentropy(), metrics=['accur
```

```
In [ ]: 1 model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 148, 148, 16)	448
max_pooling2d (MaxPooling2D)	(None, 74, 74, 16)	0
conv2d_1 (Conv2D)	(None, 36, 36, 32)	4640
max_pooling2d_1 (MaxPooling2D)	(None, 36, 36, 32)	0
conv2d_2 (Conv2D)	(None, 17, 17, 64)	18496
max_pooling2d_2 (MaxPooling2D)	(None, 17, 17, 64)	0
flatten (Flatten)	(None, 18496)	0
dense (Dense)	(None, 256)	4735232
dense_1 (Dense)	(None, 128)	32896
dense_2 (Dense)	(None, 64)	8256
dense_3 (Dense)	(None, 5)	325

```
=====
Total params: 4,800,293
Trainable params: 4,800,293
Non-trainable params: 0
```


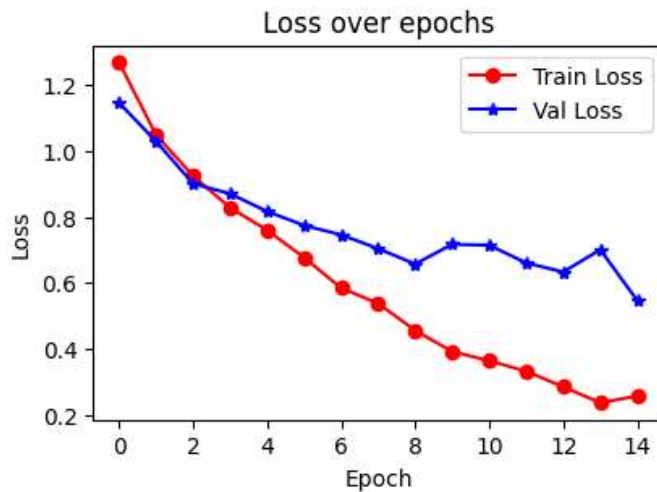
```
In [ ]: 1 history = model.fit(train_dataset, epochs=15, validation_data=(val_dataset))

Epoch 1/15
125/125 [=====] - 69s 411ms/step - loss: 1.2671 - accuracy: 0.4565 - val_
loss: 1.1449 - val_accuracy: 0.5250
Epoch 2/15
125/125 [=====] - 50s 399ms/step - loss: 1.0477 - accuracy: 0.5885 - val_
loss: 1.0282 - val_accuracy: 0.6100
Epoch 3/15
125/125 [=====] - 48s 387ms/step - loss: 0.9237 - accuracy: 0.6470 - val_
loss: 0.9004 - val_accuracy: 0.6420
Epoch 4/15
125/125 [=====] - 48s 386ms/step - loss: 0.8279 - accuracy: 0.6775 - val_
loss: 0.8699 - val_accuracy: 0.6850
Epoch 5/15
125/125 [=====] - 49s 390ms/step - loss: 0.7596 - accuracy: 0.7138 - val_
loss: 0.8166 - val_accuracy: 0.6950
Epoch 6/15
125/125 [=====] - 50s 398ms/step - loss: 0.6770 - accuracy: 0.7347 - val_
loss: 0.7740 - val_accuracy: 0.7160
Epoch 7/15
125/125 [=====] - 53s 422ms/step - loss: 0.5846 - accuracy: 0.7760 - val_
loss: 0.7458 - val_accuracy: 0.7430
Epoch 8/15
125/125 [=====] - 49s 390ms/step - loss: 0.5382 - accuracy: 0.8023 - val_
loss: 0.7037 - val_accuracy: 0.7670
Epoch 9/15
125/125 [=====] - 48s 387ms/step - loss: 0.4564 - accuracy: 0.8330 - val_
loss: 0.6574 - val_accuracy: 0.7850
Epoch 10/15
125/125 [=====] - 49s 395ms/step - loss: 0.3924 - accuracy: 0.8593 - val_
loss: 0.7169 - val_accuracy: 0.7680
Epoch 11/15
125/125 [=====] - 52s 418ms/step - loss: 0.3647 - accuracy: 0.8748 - val_
loss: 0.7141 - val_accuracy: 0.7810
Epoch 12/15
125/125 [=====] - 48s 388ms/step - loss: 0.3328 - accuracy: 0.8815 - val_
loss: 0.6606 - val_accuracy: 0.8000
Epoch 13/15
125/125 [=====] - 49s 391ms/step - loss: 0.2856 - accuracy: 0.9035 - val_
loss: 0.6334 - val_accuracy: 0.8220
Epoch 14/15
125/125 [=====] - 50s 399ms/step - loss: 0.2379 - accuracy: 0.9160 - val_
loss: 0.7007 - val_accuracy: 0.8260
Epoch 15/15
125/125 [=====] - 48s 389ms/step - loss: 0.2576 - accuracy: 0.9120 - val_
loss: 0.5470 - val_accuracy: 0.8450
```

```

In [ ]: 1 loss_df = pd.DataFrame(history.history)
2
3 def plot_predictions(data=loss_df):
4     fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(10, 3))
5
6     ax1.plot(loss_df['loss'], color='red', marker='o', label='Train Loss')
7     ax1.plot(loss_df['val_loss'], color='blue', marker='*', label='Val Loss')
8
9     ax1.set_title('Loss over epochs')
10    ax1.set_xlabel('Epoch')
11    ax1.set_ylabel('Loss')
12    ax1.legend()
13
14    ax2.plot(loss_df['accuracy'], color='red', marker='o', label='Train Accuracy')
15    ax2.plot(loss_df['val_accuracy'], color='blue', marker='*', label='Val Accuracy')
16
17    ax2.set_title('Accuracy over epochs')
18    ax2.set_xlabel('Epoch')
19    ax2.set_ylabel('Accuracy')
20    ax2.legend()
21    plot_predictions(loss_df)

```



making predictions

```

In [ ]: 1 predictions = np.argmax(model.predict(val_dataset), axis=1)
2 y_true = val_dataset.labels

```

32/32 [=====] - 6s 172ms/step

```

In [ ]: 1 print(accuracy_score(predictions, y_true))

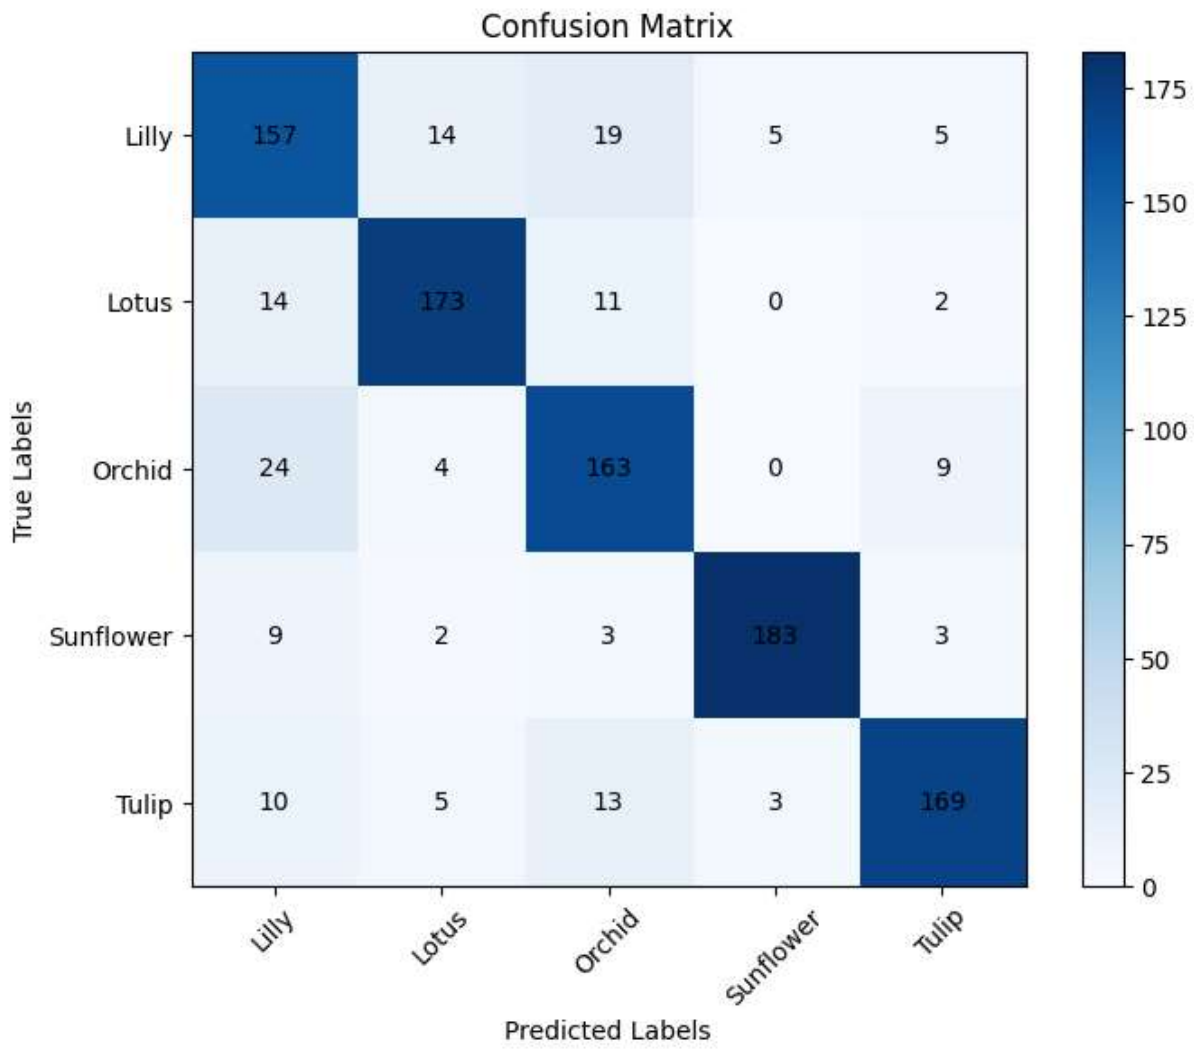
```

0.845

```

In [ ]: 1 def plot_confusion_matrix(y_true, predictions, class_names):
2
3         cm = confusion_matrix(y_true, predictions)
4         plt.figure(figsize=(8, 6))
5         heatmap = plt.imshow(cm, cmap='Blues')
6
7         # Set axis labels and title
8         plt.xlabel('Predicted Labels')
9         plt.ylabel('True Labels')
10        plt.title('Confusion Matrix')
11
12        # Set xticks and yticks with class names
13        tick_labels = class_names
14        plt.xticks(ticks=np.arange(len(class_names)), labels=tick_labels, rotation=45)
15        plt.yticks(ticks=np.arange(len(class_names)), labels=tick_labels)
16
17        # Add numbers to the heatmap cells
18        for i in range(len(class_names)):
19            for j in range(len(class_names)):
20                plt.text(j, i, str(cm[i, j]), ha='center', va='center', color='black')
21
22        plt.colorbar(heatmap)
23        plt.show()
24    plot_confusion_matrix(y_true, predictions, class_names)

```




```

In [ ]: 1 def plot_random_image(model, val_data, classes):
2
3     images = []
4     labels = []
5     for _ in range(len(val_data)):
6         batch_images, batch_labels = next(val_data)
7         images.extend(batch_images)
8         labels.extend(batch_labels)
9
10    # Shuffle the images and labels together
11    combined = list(zip(images, labels))
12    random.shuffle(combined)
13    images, labels = zip(*combined)
14    labels = np.argmax(labels, axis=1)
15    plt.figure(figsize=(12, 6))
16    for i in range(6):
17        ax = plt.subplot(2, 3, i + 1)
18        rand_index = random.choice(range(len(images)))
19        target_image = images[rand_index]
20        pred_probs = model.predict(tf.expand_dims(target_image, axis=0), verbose=0)
21        pred_label = classes[pred_probs.argmax()]
22        true_label = classes[labels[rand_index]]
23
24        plt.imshow(target_image)
25
26        if pred_label == true_label:
27            color = "green"
28        else:
29            color = "red"
30
31        plt.title("Pred: {} {:.2f}% (True: {})".format(pred_label,
32                                                       100 * tf.reduce_max(pred_probs),
33                                                       true_label),
34                color=color, fontsize=10)
35
36    plt.tight_layout()
37    plot_random_image(model, val_dataset, class_names)
38    plt.show()

```

