


Apple Leaf Disease Detection



```
import numpy as np
import pandas as pd
import os

base_path = "/kaggle/input/leaf-disease-images/leaves/Apple"
categories = ["Apple_Black_rot", "Apple_healthy", "Apple_rust", "Apple_scab"]

image_paths = []
labels = []

for category in categories:
    category_path = os.path.join(base_path, category)
    for image_name in os.listdir(category_path):
        image_path = os.path.join(category_path, image_name)
        image_paths.append(image_path)
        labels.append(category)

df = pd.DataFrame({
    "image_path": image_paths,
    "label": labels
})

df.head()
```

	image_path	label
0	/kaggle/input/leaf-disease-images/leaves/Apple...	Apple_Black_rot
1	/kaggle/input/leaf-disease-images/leaves/Apple...	Apple_Black_rot

```

2 /kaggle/input/leaf-disease-images/leaves/Apple... Apple_Black_rot
3 /kaggle/input/leaf-disease-images/leaves/Apple... Apple_Black_rot
4 /kaggle/input/leaf-disease-images/leaves/Apple... Apple_Black_rot

```

```
df.tail()
```

```

                                image_path      label
3159 /kaggle/input/leaf-disease-images/leaves/Apple... Apple_scab
3160 /kaggle/input/leaf-disease-images/leaves/Apple... Apple_scab
3161 /kaggle/input/leaf-disease-images/leaves/Apple... Apple_scab
3162 /kaggle/input/leaf-disease-images/leaves/Apple... Apple_scab
3163 /kaggle/input/leaf-disease-images/leaves/Apple... Apple_scab

```

```
df.shape
```

```
(3164, 2)
```

```
df.columns
```

```
Index(['image_path', 'label'], dtype='object')
```

```
df.duplicated().sum()
```

```
0
```

```
df.isnull().sum()
```

```

image_path    0
label         0
dtype: int64

```

```
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3164 entries, 0 to 3163
Data columns (total 2 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   image_path  3164 non-null   object
 1   label       3164 non-null   object
dtypes: object(2)
memory usage: 49.6+ KB

```

```
df['label'].unique()
```

```

array(['Apple_Black_rot', 'Apple_healthy', 'Apple_rust', 'Apple_scab'],
      dtype=object)

```

```
df['label'].value_counts()
```

```

label
Apple_healthy    1640
Apple_scab       629

```

```
Apple_Black_rot      620
Apple_rust            275
Name: count, dtype: int64
```

```
import seaborn as sns
import matplotlib.pyplot as plt
```

```
sns.set_style("whitegrid")
```

```
fig, ax = plt.subplots(figsize=(8, 6))
sns.countplot(data=df, x="label", palette="viridis", ax=ax)
```

```
ax.set_title("Distribution of Disease Types", fontsize=14, fontweight='bold')
ax.set_xlabel("Tumor Type", fontsize=12)
ax.set_ylabel("Count", fontsize=12)
```

```
for p in ax.patches:
    ax.annotate(f'{int(p.get_height())}',
                (p.get_x() + p.get_width() / 2., p.get_height()),
                ha='center', va='bottom', fontsize=11, color='black',
                xytext=(0, 5), textcoords='offset points')
```

```
plt.xticks(rotation=-45)
plt.show()
```


```
label_counts = df["label"].value_counts()
```

```
fig, ax = plt.subplots(figsize=(20, 8))
colors = sns.color_palette("viridis", len(label_counts))
```

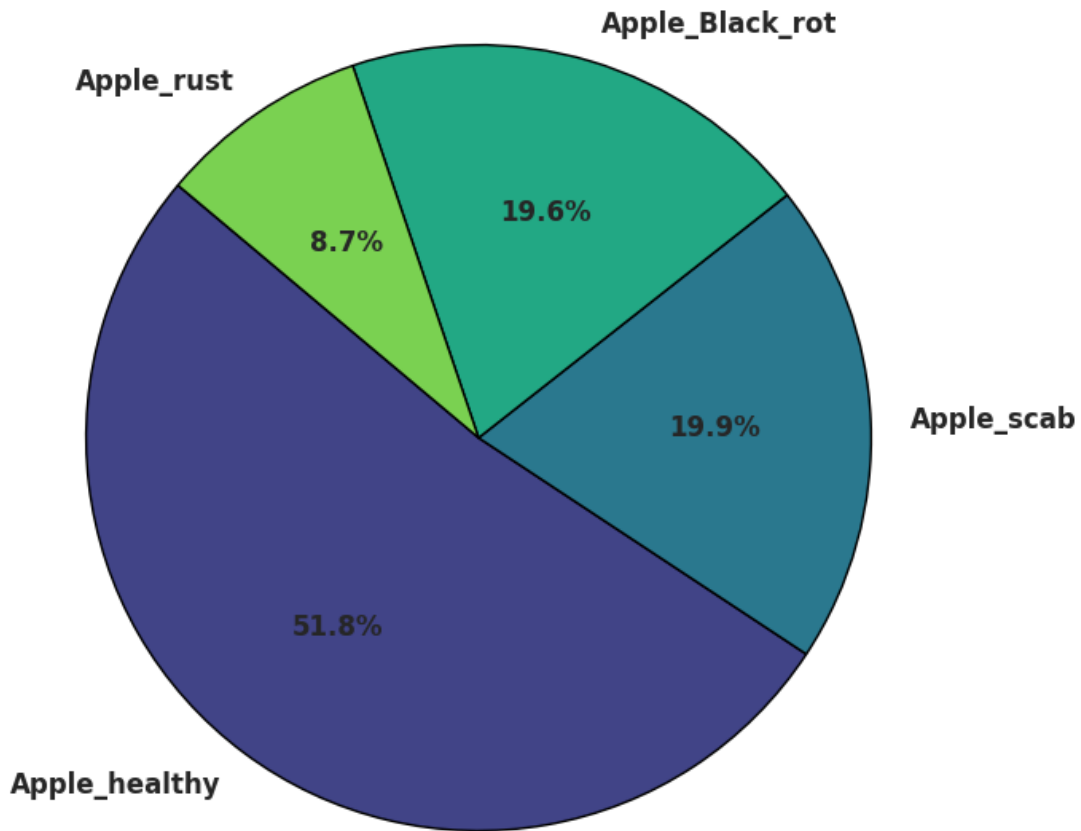
```
ax.pie(label_counts, labels=label_counts.index, autopct='%1.1f%%',
        startangle=140, colors=colors, textprops={'fontsize': 12, 'weight':
        'bold'},
        wedgeprops={'edgecolor': 'black', 'linewidth': 1})
```

```
ax.set_title("Distribution of Disease Types - Pie Chart", fontsize=14,
fontweight='bold')
```

```
plt.show()
```



Distribution of Disease Types - Pie Chart



```
import cv2

num_images = 5

plt.figure(figsize=(15, 12))

for i, category in enumerate(categories):
    category_images = df[df['label'] ==
category]['image_path'].iloc[:num_images]

    for j, img_path in enumerate(category_images):

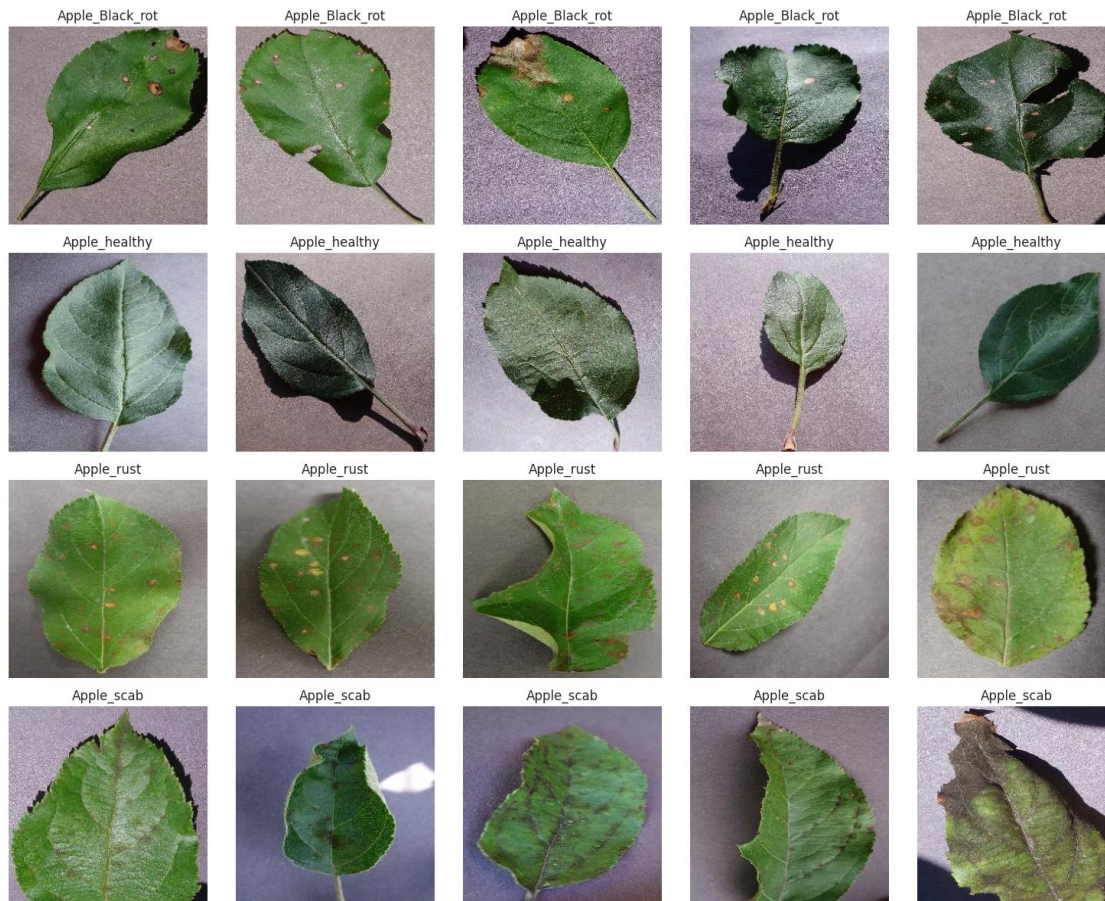
        img = cv2.imread(img_path)
        img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)

        plt.subplot(len(categories), num_images, i * num_images + j + 1)
        plt.imshow(img)
        plt.axis('off')
```

```
plt.title(category)
```

```
plt.tight_layout()
```

```
plt.show()
```



```
from sklearn.preprocessing import LabelEncoder
label_encoder = LabelEncoder()
df['category_encoded'] = label_encoder.fit_transform(df['label'])

df = df[['image_path', 'category_encoded']]

from sklearn.utils import resample

max_count = df['category_encoded'].value_counts().max()

dfs = []
for category in df['category_encoded'].unique():
    class_subset = df[df['category_encoded'] == category]
    class_upsampled = resample(class_subset,
                              replace=True,
                              n_samples=max_count,
                              random_state=42)
    dfs.append(class_upsampled)
```

```
df_balanced = pd.concat(dfs).sample(frac=1,  
random_state=42).reset_index(drop=True)
```

```
df_balanced['category_encoded'].value_counts()
```

```
category_encoded
```

```
1    1640
```

```
3    1640
```

```
2    1640
```

```
0    1640
```

```
Name: count, dtype: int64
```

```
df_resampled = df_balanced
```

```
df_resampled['category_encoded'] =  
df_resampled['category_encoded'].astype(str)
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.metrics import confusion_matrix, classification_report
```

```
import tensorflow as tf
```

```
from tensorflow import keras
```

```
from tensorflow.keras.models import Sequential
```

```
from tensorflow.keras.optimizers import Adam
```

```
from tensorflow.keras.preprocessing.image import ImageDataGenerator
```

```
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense,  
Activation, Dropout, BatchNormalization
```

```
from tensorflow.keras import regularizers
```

```
import warnings
```

```
warnings.filterwarnings("ignore")
```

```
print ('check')
```

```
2025-06-10 12:32:53.447914: E
```

```
external/local_xla/xla/stream_executor/cuda/cuda_fft.cc:477] Unable to  
register cuFFT factory: Attempting to register factory for plugin cuFFT when  
one has already been registered
```

```
WARNING: All log messages before absl::InitializeLog() is called are written  
to STDERR
```

```
E0000 00:00:1749558773.671901      35 cuda_dnn.cc:8310] Unable to register  
cuDNN factory: Attempting to register factory for plugin cuDNN when one has  
already been registered
```

```
E0000 00:00:1749558773.738947      35 cuda_blas.cc:1418] Unable to register  
cuBLAS factory: Attempting to register factory for plugin cuBLAS when one has  
already been registered
```

```
check
```



```

train_df_new, temp_df_new = train_test_split(
    df_resampled,
    train_size=0.8,
    shuffle=True,
    random_state=42,
    stratify=df_resampled['category_encoded']
)

valid_df_new, test_df_new = train_test_split(
    temp_df_new,
    test_size=0.5,
    shuffle=True,
    random_state=42,
    stratify=temp_df_new['category_encoded']
)

from tensorflow.keras.preprocessing.image import ImageDataGenerator

batch_size = 16
img_size = (224, 224)
channels = 3
img_shape = (img_size[0], img_size[1], channels)

tr_gen = ImageDataGenerator(
    rescale=1./255
)

ts_gen = ImageDataGenerator(rescale=1./255)

train_gen_new = tr_gen.flow_from_dataframe(
    train_df_new,
    x_col='image_path',
    y_col='category_encoded',
    target_size=img_size,
    class_mode='sparse',
    color_mode='rgb',
    shuffle=True,
    batch_size=batch_size
)

valid_gen_new = ts_gen.flow_from_dataframe(
    valid_df_new,
    x_col='image_path',
    y_col='category_encoded',
    target_size=img_size,
    class_mode='sparse',
    color_mode='rgb',
    shuffle=True,
    batch_size=batch_size
)

```



```
)

test_gen_new = ts_gen.flow_from_dataframe(
    test_df_new,
    x_col='image_path',
    y_col='category_encoded',
    target_size=img_size,
    class_mode='sparse',
    color_mode='rgb',
    shuffle=False,
    batch_size=batch_size
)
```

Found 5248 validated image filenames belonging to 4 classes.

Found 656 validated image filenames belonging to 4 classes.

Found 656 validated image filenames belonging to 4 classes.

```
print("Num GPUs Available: ", len(tf.config.list_physical_devices('GPU')))
```

Num GPUs Available: 2

```
gpus = tf.config.list_physical_devices('GPU')
if gpus:
    try:
        for gpu in gpus:
            tf.config.experimental.set_memory_growth(gpu, True)
        print("GPU is set for TensorFlow")
    except RuntimeError as e:
        print(e)
```

GPU is set for TensorFlow

```
from tensorflow.keras import layers, models
```

```
num_classes = 4
```

```
import tensorflow as tf
from tensorflow.keras import layers, models
```

```
class ContinuousLayer(layers.Layer):
    def __init__(self, kernel_size=5, num_basis=10, output_channels=16,
**kwargs):
        super(ContinuousLayer, self).__init__(**kwargs)
        self.kernel_size = kernel_size
        self.num_basis = num_basis
        self.output_channels = output_channels
        self.centers = self.add_weight(
            name='centers',
            shape=(num_basis, 2),
            initializer='random_normal',
            trainable=True
```

```

    )
    self.widths = self.add_weight(
        name='widths',
        shape=(num_basis,),
        initializer='ones',
        trainable=True,
        constraint=tf.keras.constraints.NonNeg()
    )
    self.kernel_weights = self.add_weight(
        name='kernel_weights',
        shape=(kernel_size, kernel_size, 32, output_channels), # Updated
to 32 input channels
        initializer='glorot_normal',
        trainable=True
    )

    def call(self, inputs):
        height, width = img_size
        x = tf.range(0, height, 1.0)
        y = tf.range(0, width, 1.0)
        x_grid, y_grid = tf.meshgrid(x, y)
        grid = tf.stack([x_grid, y_grid], axis=-1)

        basis = []
        for i in range(self.num_basis):
            center = self.centers[i]
            width = self.widths[i]
            dist = tf.reduce_sum(((grid - center) / width) ** 2, axis=-1)
            basis_i = tf.exp(-dist)
            basis.append(basis_i)
        basis = tf.stack(basis, axis=-1)

        basis_weights = tf.reduce_mean(basis, axis=[0, 1])
        basis_weights = tf.nn.softmax(basis_weights)
        basis_weights = basis_weights[:, tf.newaxis, tf.newaxis, tf.newaxis,
tf.newaxis]

        modulated_kernel = self.kernel_weights * tf.reduce_sum(basis_weights,
axis=0)

        output = tf.nn.conv2d(
            inputs,
            modulated_kernel,
            strides=[1, 1, 1, 1],
            padding='SAME'
        )

    return output

```

```

    def compute_output_shape(self, input_shape):
        return (input_shape[0], input_shape[1], input_shape[2],
self.output_channels)

    def smoothness_penalty(self):
        grad_x = tf.reduce_mean(tf.square(self.kernel_weights[1:, :, :, :] -
self.kernel_weights[:-1, :, :, :]))
        grad_y = tf.reduce_mean(tf.square(self.kernel_weights[:, 1:, :, :] -
self.kernel_weights[:, :-1, :, :]))
        return grad_x + grad_y

class VariationalLoss(tf.keras.losses.Loss):
    def __init__(self, model, lambda1=0.01, lambda2=1.0):
        super(VariationalLoss, self).__init__()
        self.model = model
        self.lambda1 = lambda1
        self.lambda2 = lambda2
        self.sce = tf.keras.losses.SparseCategoricalCrossentropy()

    def call(self, y_true, y_pred):
        smoothness_penalty = 0
        for layer in self.model.layers:
            if isinstance(layer, ContinuousLayer):
                smoothness_penalty += layer.smoothness_penalty()
        prediction_loss = self.sce(y_true, y_pred)
        return self.lambda2 * prediction_loss + self.lambda1 *
smoothness_penalty

def build_continuous_model():
    inputs = layers.Input(shape=img_shape)
    x = layers.Conv2D(filters=32, kernel_size=3, padding='same',
activation='relu')(inputs)
    x = ContinuousLayer(kernel_size=5, num_basis=10, output_channels=16)(x)
    x = layers.Activation('relu')(x)
    x = layers.MaxPooling2D(pool_size=(2, 2))(x)
    x = layers.Flatten()(x)
    x = layers.Dense(128, activation='relu')(x)
    x = layers.Dropout(0.5)(x)
    outputs = layers.Dense(num_classes, activation='softmax')(x)
    model = models.Model(inputs, outputs)
    return model

model = build_continuous_model()

model.compile(
    optimizer='adam',
    loss=VariationalLoss(model=model, lambda1=0.01, lambda2=1.0),
    metrics=['accuracy']
)

```

```
history = model.fit(
    train_gen_new,
    validation_data=valid_gen_new,
    epochs=10,
    verbose=1
)
```

Epoch 1/10

WARNING: All log messages before absl::InitializeLog() is called are written to STDERR

```
I0000 00:00:1749558926.317944      84 service.cc:148] XLA service
0x7e4790015780 initialized for platform CUDA (this does not guarantee that
XLA will be used). Devices:
I0000 00:00:1749558926.318897      84 service.cc:156]   StreamExecutor device
(0): Tesla T4, Compute Capability 7.5
I0000 00:00:1749558926.318931      84 service.cc:156]   StreamExecutor device
(1): Tesla T4, Compute Capability 7.5
I0000 00:00:1749558926.794642      84 cuda_dnn.cc:529] Loaded cuDNN version
90300
```

3/328 ————— 20s 63ms/step - accuracy: 0.2812 - loss: 8.2690

```
I0000 00:00:1749558934.082973      84 device_compiler.h:188] Compiled cluster
using XLA! This line is logged at most once for the lifetime of the process.
```

328/328 ————— 33s 66ms/step - accuracy: 0.5818 - loss: 1.7665
- val_accuracy: 0.9131 - val_loss: 0.2266

Epoch 2/10

328/328 ————— 14s 43ms/step - accuracy: 0.9091 - loss: 0.2781
- val_accuracy: 0.9466 - val_loss: 0.1548

Epoch 3/10

328/328 ————— 13s 40ms/step - accuracy: 0.9532 - loss: 0.1428
- val_accuracy: 0.9558 - val_loss: 0.1355

Epoch 4/10

328/328 ————— 13s 41ms/step - accuracy: 0.9713 - loss: 0.0845
- val_accuracy: 0.9573 - val_loss: 0.1367

Epoch 5/10

328/328 ————— 13s 40ms/step - accuracy: 0.9734 - loss: 0.0738
- val_accuracy: 0.9573 - val_loss: 0.1654

Epoch 6/10

328/328 ————— 13s 40ms/step - accuracy: 0.9807 - loss: 0.0772
- val_accuracy: 0.9665 - val_loss: 0.1116

Epoch 7/10

328/328 ————— 14s 41ms/step - accuracy: 0.9889 - loss: 0.0395
- val_accuracy: 0.9619 - val_loss: 0.1281

Epoch 8/10

328/328 ————— 14s 41ms/step - accuracy: 0.9763 - loss: 0.0641
- val_accuracy: 0.9695 - val_loss: 0.0987

Epoch 9/10

```

328/328 -----14s 42ms/step - accuracy: 0.9717 - loss: 0.1013
- val_accuracy: 0.9726 - val_loss: 0.1221
Epoch 10/10
328/328 -----14s 41ms/step - accuracy: 0.9836 - loss: 0.0656
- val_accuracy: 0.9695 - val_loss: 0.1543

```

```
model.summary()
```

```
Model: "functional_1"
```

Layer (type) Param #	Output Shape	
input_layer_2 (InputLayer) 0	(None, 224, 224, 3)	
conv2d_2 (Conv2D) 896	(None, 224, 224, 32)	
continuous_layer_2 (ContinuousLayer) 12,830	(None, 224, 224, 16)	
activation_2 (Activation) 0	(None, 224, 224, 16)	
max_pooling2d_2 (MaxPooling2D) 0	(None, 112, 112, 16)	
flatten_2 (Flatten) 0	(None, 200704)	
dense_3 (Dense) 25,690,240	(None, 128)	
dropout_2 (Dropout) 0	(None, 128)	
dense_4 (Dense) 516	(None, 4)	

Total params: 77,113,448 (294.16 MB)

Trainable params: 25,704,482 (98.05 MB)

Non-trainable params: 0 (0.00 B)

Optimizer params: 51,408,966 (196.11 MB)

```
test_loss, test_accuracy = model.evaluate(test_gen_new)
print(f"Test Loss: {test_loss:.4f}, Test Accuracy: {test_accuracy:.4f}")
```

```
41/41 ————— 2s 43ms/step - accuracy: 0.9861 - loss: 0.0634
Test Loss: 0.0620, Test Accuracy: 0.9863
```

```
plt.figure(figsize=(12, 4))
plt.subplot(1, 2, 1)
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Loss Over Epochs')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Accuracy Over Epochs')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
plt.show()
```



```
test_gen_new.reset()
y_pred = model.predict(test_gen_new)
```

```

y_pred_classes = np.argmax(y_pred, axis=1)
y_true = test_gen_new.classes


41/41 ————— 2s 31ms/step

cm = confusion_matrix(y_true, y_pred_classes)

class_names = list(test_gen_new.class_indices.keys())

plt.figure(figsize=(10, 8))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=class_names,
yticklabels=class_names)
plt.title('Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.show()

```



```

from sklearn.metrics import classification_report

print("\nClassification Report:")
print(classification_report(y_true, y_pred_classes))

```


Classification Report:

	precision	recall	f1-score	support
0	0.98	0.99	0.99	164
1	1.00	0.96	0.98	164
2	1.00	1.00	1.00	164
3	0.96	0.99	0.98	164
accuracy			0.99	656
macro avg	0.99	0.99	0.99	656
weighted avg	0.99	0.99	0.99	656