

# Predicting Cervical Cancer

## STEP 1: IMPORTING LIBRARIES AND DATASETS

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
sns.set()
```

In [74]:

```
df = pd.read_csv('cervical_cancer.csv')

# display all columns
pd.set_option('display.max_columns', None)
df
```

In [75]:

Out[75]:

	Age	Number of sexual partners	First sexual intercourse	Num of pregnancies	Smokes	Smokes (years)	Smokes (packs/year)	Hormonal Contraceptives	Hormonal Contraceptives (years)	IUD	IUD (years)	STDs	STDs (number)	ST
0	18	4.0	15.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
1	15	1.0	14.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
2	34	1.0	?	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
3	52	5.0	16.0	4.0	1.0	37.0	37.0	1.0	3.0	0.0	0.0	0.0	0.0	0.0
4	46	3.0	21.0	4.0	0.0	0.0	0.0	1.0	15.0	0.0	0.0	0.0	0.0	0.0
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
853	34	3.0	18.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
854	32	2.0	19.0	1.0	0.0	0.0	0.0	1.0	8.0	0.0	0.0	0.0	0.0	0.0
855	25	2.0	17.0	0.0	0.0	0.0	0.0	1.0	0.08	0.0	0.0	0.0	0.0	0.0
856	33	2.0	24.0	2.0	0.0	0.0	0.0	1.0	0.08	0.0	0.0	0.0	0.0	0.0
857	29	2.0	20.0	1.0	0.0	0.0	0.0	1.0	0.5	0.0	0.0	0.0	0.0	0.0

858 rows × 36 columns

this is a binary class classification

## step 2: Data Analysis

```
df.describe()
```

In [76]:

Out[76]:

	Age	STDs: Number of diagnosis	Dx:Cancer	Dx:CIN	Dx:HPV	Dx	Hinselmann	Schiller	Cit
<b>count</b>	858.000000	858.000000	858.000000	858.000000	858.000000	858.000000	858.000000	858.000000	858.000000
<b>mean</b>	26.820513	0.087413	0.020979	0.010490	0.020979	0.027972	0.040793	0.086247	0.051247
<b>std</b>	8.497948	0.302545	0.143398	0.101939	0.143398	0.164989	0.197925	0.280892	0.220747
<b>min</b>	13.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
<b>25%</b>	20.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
<b>50%</b>	25.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
<b>75%</b>	32.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
<b>max</b>	84.000000	3.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000

In [77]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 858 entries, 0 to 857
```

```
Data columns (total 36 columns):
```

#	Column	Non-Null Count	Dtype
---	-----	-----	-----
0	Age	858 non-null	int64
1	Number of sexual partners	858 non-null	object
2	First sexual intercourse	858 non-null	object
3	Num of pregnancies	858 non-null	object
4	Smokes	858 non-null	object
5	Smokes (years)	858 non-null	object
6	Smokes (packs/year)	858 non-null	object
7	Hormonal Contraceptives	858 non-null	object
8	Hormonal Contraceptives (years)	858 non-null	object
9	IUD	858 non-null	object
10	IUD (years)	858 non-null	object
11	STDs	858 non-null	object
12	STDs (number)	858 non-null	object
13	STDs:condylomatosis	858 non-null	object
14	STDs:cervical condylomatosis	858 non-null	object
15	STDs:vaginal condylomatosis	858 non-null	object
16	STDs:vulvo-perineal condylomatosis	858 non-null	object
17	STDs:syphilis	858 non-null	object
18	STDs:pelvic inflammatory disease	858 non-null	object
19	STDs:genital herpes	858 non-null	object
20	STDs:molluscum contagiosum	858 non-null	object
21	STDs:AIDS	858 non-null	object
22	STDs:HIV	858 non-null	object

23	STDs:Hepatitis B	858 non-null	object
24	STDs:HPV	858 non-null	object
25	STDs: Number of diagnosis	858 non-null	int64
26	STDs: Time since first diagnosis	858 non-null	object
27	STDs: Time since last diagnosis	858 non-null	object
28	Dx:Cancer	858 non-null	int64
29	Dx:CIN	858 non-null	int64
30	Dx:HPV	858 non-null	int64
31	Dx	858 non-null	int64
32	Hinselmann	858 non-null	int64
33	Schiller	858 non-null	int64
34	Citology	858 non-null	int64
35	Biopsy	858 non-null	int64

dtypes: int64(10), object(26)  
memory usage: 241.4+ KB

In [78]:

```
# checking the null values
df.isnull().sum()
```

Out[78]:

Age	0
Number of sexual partners	0
First sexual intercourse	0
Num of pregnancies	0
Smokes	0
Smokes (years)	0
Smokes (packs/year)	0
Hormonal Contraceptives	0
Hormonal Contraceptives (years)	0
IUD	0
IUD (years)	0
STDs	0
STDs (number)	0
STDs:condylomatosis	0
STDs:cervical condylomatosis	0
STDs:vaginal condylomatosis	0
STDs:vulvo-perineal condylomatosis	0
STDs:syphilis	0
STDs:pelvic inflammatory disease	0
STDs:genital herpes	0
STDs:molluscum contagiosum	0
STDs:AIDS	0
STDs:HIV	0
STDs:Hepatitis B	0
STDs:HPV	0
STDs: Number of diagnosis	0
STDs: Time since first diagnosis	0

```
STDs: Time since last diagnosis      0
Dx:Cancer                           0
Dx:CIN                              0
Dx:HPV                              0
Dx                                  0
Hinselmann                          0
Schiller                             0
Citology                             0
Biopsy                              0
dtype: int64
no null values
```

In [79]:

```
# replacing '?' with Null
df = df.replace(to_replace='?', value=np.nan)
df
```

Out[79]:

	Age	Number of sexual partners	First sexual intercourse	Num of pregnancies	Smokes	Smokes (years)	Smokes (packs/year)	Hormonal Contraceptives	Hormonal Contraceptives (years)	IUD	IUD (years)	STDs	STDs (number)	STDs (number)
0	18	4.0	15.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
1	15	1.0	14.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
2	34	1.0	NaN	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
3	52	5.0	16.0	4.0	1.0	37.0	37.0	1.0	3.0	0.0	0.0	0.0	0.0	0.0
4	46	3.0	21.0	4.0	0.0	0.0	0.0	1.0	15.0	0.0	0.0	0.0	0.0	0.0
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
853	34	3.0	18.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
854	32	2.0	19.0	1.0	0.0	0.0	0.0	1.0	8.0	0.0	0.0	0.0	0.0	0.0
855	25	2.0	17.0	0.0	0.0	0.0	0.0	1.0	0.08	0.0	0.0	0.0	0.0	0.0
856	33	2.0	24.0	2.0	0.0	0.0	0.0	1.0	0.08	0.0	0.0	0.0	0.0	0.0
857	29	2.0	20.0	1.0	0.0	0.0	0.0	1.0	0.5	0.0	0.0	0.0	0.0	0.0

858 rows × 36 columns

In [80]:

```
# checking the number of NAN in the dataset
df.isnull().sum()
```

Out[80]:

Age 0

Number of sexual partners	26
First sexual intercourse	7
Num of pregnancies	56
Smokes	13
Smokes (years)	13
Smokes (packs/year)	13
Hormonal Contraceptives	108
Hormonal Contraceptives (years)	108
IUD	117
IUD (years)	117
STDs	105
STDs (number)	105
STDs:condylomatosis	105
STDs:cervical condylomatosis	105
STDs:vaginal condylomatosis	105
STDs:vulvo-perineal condylomatosis	105
STDs:syphilis	105
STDs:pelvic inflammatory disease	105
STDs:genital herpes	105
STDs:molluscum contagiosum	105
STDs:AIDS	105
STDs:HIV	105
STDs:Hepatitis B	105
STDs:HPV	105
STDs: Number of diagnosis	0
STDs: Time since first diagnosis	787
STDs: Time since last diagnosis	787
Dx:Cancer	0
Dx:CIN	0
Dx:HPV	0
Dx	0
Hinselmann	0
Schiller	0
Citology	0
Biopsy	0
dtype: int64	

In [81]:

```
# plotting a heat map to visualize the number of nan in the data
plt.figure(figsize=(20, 20), facecolor="white")
sns.heatmap(df.isnull(), yticklabels=False)
```

Out[81]:

<AxesSubplot:>



There are a lot of null values in the columns "STDs: Time since first diagnosis" and "STDs: Time since last diagnosis" This columns will need to be dropped.

In [82]:

```
df = df.drop(columns=['STDs: Time since first diagnosis', 'STDs: Time since
last diagnosis'], axis=1)
df
```

Out[82]:

	Age	Number of sexual partners	First sexual intercourse	Num of pregnancies	Smokes	Smokes (years)	Smokes (packs/year)	Hormonal Contraceptives	Hormonal Contraceptives (years)	IUD	IUD (years)	STDs	STD (number)
0	18	4.0	15.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
1	15	1.0	14.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
2	34	1.0	NaN	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
3	52	5.0	16.0	4.0	1.0	37.0	37.0	1.0	3.0	0.0	0.0	0.0	0.0
4	46	3.0	21.0	4.0	0.0	0.0	0.0	1.0	15.0	0.0	0.0	0.0	0.0
...	...	...	...	...	...	...	...	...	...	...	...	...	...
853	34	3.0	18.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
854	32	2.0	19.0	1.0	0.0	0.0	0.0	1.0	8.0	0.0	0.0	0.0	0.0
855	25	2.0	17.0	0.0	0.0	0.0	0.0	1.0	0.08	0.0	0.0	0.0	0.0
856	33	2.0	24.0	2.0	0.0	0.0	0.0	1.0	0.08	0.0	0.0	0.0	0.0
857	29	2.0	20.0	1.0	0.0	0.0	0.0	1.0	0.5	0.0	0.0	0.0	0.0

858 rows × 34 columns

my dataframe contains objects as data type which is so generic, we need to convert the columns to numeric data type, which handles float, int and even strings which are in form of numbers. this will aid to perform statistical analysis.

In [83]:

```
df = df.apply(pd.to_numeric)
df.info()
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 858 entries, 0 to 857
Data columns (total 34 columns):
```

#	Column	Non-Null Count	Dtype
---	-----	-----	-----
0	Age	858 non-null	int64
1	Number of sexual partners	832 non-null	float64
2	First sexual intercourse	851 non-null	float64
3	Num of pregnancies	802 non-null	float64
4	Smokes	845 non-null	float64

```

5  Smokes (years)                845 non-null    float64
6  Smokes (packs/year)          845 non-null    float64
7  Hormonal Contraceptives      750 non-null    float64
8  Hormonal Contraceptives (years) 750 non-null    float64
9  IUD                          741 non-null    float64
10 IUD (years)                  741 non-null    float64
11 STDs                         753 non-null    float64
12 STDs (number)                753 non-null    float64
13 STDs:condylomatosis          753 non-null    float64
14 STDs:cervical condylomatosis 753 non-null    float64
15 STDs:vaginal condylomatosis  753 non-null    float64
16 STDs:vulvo-perineal condylomatosis 753 non-null    float64
17 STDs:syphilis                753 non-null    float64
18 STDs:pelvic inflammatory disease 753 non-null    float64
19 STDs:genital herpes          753 non-null    float64
20 STDs:molluscum contagiosum   753 non-null    float64
21 STDs:AIDS                    753 non-null    float64
22 STDs:HIV                     753 non-null    float64
23 STDs:Hepatitis B            753 non-null    float64
24 STDs:HPV                     753 non-null    float64
25 STDs: Number of diagnosis    858 non-null    int64
26 Dx:Cancer                    858 non-null    int64
27 Dx:CIN                       858 non-null    int64
28 Dx:HPV                       858 non-null    int64
29 Dx                           858 non-null    int64
30 Hinselmann                   858 non-null    int64
31 Schiller                     858 non-null    int64
32 Citology                     858 non-null    int64
33 Biopsy                       858 non-null    int64

```

dtypes: float64(24), int64(10)

memory usage: 228.0 KB

In [84]:

```

#i need to replace my null values with mode
df.mode()

```

Out[84]:

	Age	Number of sexual partners	First sexual intercourse	Num of pregnancies	Smokes	Smokes (years)	Smokes (packs/year)	Hormonal Contraceptives	Hormonal Contraceptives (years)	IUD	IUD (years)	STDs	STDs (number)
0	23	2.0	15.0	1.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0

In [95]:

```

for column in df:
    df[column] = df[column].fillna(df[column].mode()[0])
df

```

Out[95]:



	Age	Number of sexual partners	First sexual intercourse	Num of pregnancies	Smokes	Smokes (years)	Smokes (packs/year)	Hormonal Contraceptives	Hormonal Contraceptives (years)	IUD	IUD (years)	STDs	STDs (number)
0	18	4.0	15.0	1.0	0.0	0.0	0.0	0.0	0.00	0.0	0.0	0.0	0.0
1	15	1.0	14.0	1.0	0.0	0.0	0.0	0.0	0.00	0.0	0.0	0.0	0.0
2	34	1.0	15.0	1.0	0.0	0.0	0.0	0.0	0.00	0.0	0.0	0.0	0.0
3	52	5.0	16.0	4.0	1.0	37.0	37.0	1.0	3.00	0.0	0.0	0.0	0.0
4	46	3.0	21.0	4.0	0.0	0.0	0.0	1.0	15.00	0.0	0.0	0.0	0.0
...	...	...	...	...	...	...	...	...	...	...	...	...	...
853	34	3.0	18.0	0.0	0.0	0.0	0.0	0.0	0.00	0.0	0.0	0.0	0.0
854	32	2.0	19.0	1.0	0.0	0.0	0.0	1.0	8.00	0.0	0.0	0.0	0.0
855	25	2.0	17.0	0.0	0.0	0.0	0.0	1.0	0.08	0.0	0.0	0.0	0.0
856	33	2.0	24.0	2.0	0.0	0.0	0.0	1.0	0.08	0.0	0.0	0.0	0.0
857	29	2.0	20.0	1.0	0.0	0.0	0.0	1.0	0.50	0.0	0.0	0.0	0.0

858 rows × 34 columns

```

df.isnull().sum()
In [96]:

Out[96]:
Age                                0
Number of sexual partners          0
First sexual intercourse            0
Num of pregnancies                 0
Smokes                            0
Smokes (years)                    0
Smokes (packs/year)               0
Hormonal Contraceptives           0
Hormonal Contraceptives (years)   0
IUD                                0
IUD (years)                       0
STDs                               0
STDs (number)                     0
STDs:condylomatosis               0
STDs:cervical condylomatosis      0
STDs:vaginal condylomatosis       0
STDs:vulvo-perineal condylomatosis 0
STDs:syphilis                     0
STDs:pelvic inflammatory disease  0

```


```
STDs:genital herpes      0
STDs:molluscum contagiosum 0
STDs:AIDS                0
STDs:HIV                 0
STDs:Hepatitis B        0
STDs:HPV                 0
STDs: Number of diagnosis 0
Dx:Cancer                0
Dx:CIN                   0
Dx:HPV                   0
Dx                       0
Hinselmann               0
Schiller                 0
Citology                 0
Biopsy                   0
dtype: int64
```

```
plt.figure(figsize=(8, 5), facecolor='white')
sns.heatmap(df.isnull(), yticklabels=False)
```

```
<AxesSubplot:>
```

In [102]:

Out[102]:



## Data visualization

























```
# splitting my training data from the target
X = df.drop('Biopsy', axis=1)
y = df['Biopsy']

plt.figure(figsize=(20,30), facecolor='white')
plotnumber = 1
```

In [109]:

In [110]:

```
for column in X:
    if plotnumber<=len(df.columns) :
        ax = plt.subplot(12,3,plotnumber)
        plt.scatter(X[column],y)
        plt.xlabel(column,fontsize=20)
        plt.ylabel('Biopsy',fontsize=20)
    plotnumber+=1
plt.tight_layout()
```



```

In [118]:
# observing the correlation of the feature columns. columns with very high
correlation should be removed
corr_matrix = df.corr()
corr_matrix

```

Out[118]:

	Age	Number of sexual partners	First sexual intercourse	Num of pregnancies	Smokes	Smokes (years)	Smokes (packs/year)	Hormonal Contraceptives	Hormonal Contraceptives (years)	IUD	IUD (years)
Age	1.000000	0.085971	0.365248	0.514977	0.057204	0.218261	0.131861	0.029201	0.298892	0.279429	0.215427
Number of sexual partners	0.085971	1.000000	-0.147280	0.082388	0.236858	0.175729	0.174968	0.004027	0.021525	0.032460	0.006252
First sexual intercourse	0.365248	-0.147280	1.000000	-0.046099	0.123017	0.058620	-0.057013	-0.009563	0.031976	-0.008826	-0.015697
Num of pregnancies	0.514977	0.082388	-0.046099	1.000000	0.077363	0.172084	0.092214	0.116944	0.221456	0.198134	0.148692
Smokes	0.057204	0.236858	-0.123017	0.077363	1.000000	0.723572	0.493843	0.004036	0.040917	-0.055115	-0.035798
Smokes (years)	0.218261	0.175729	-0.058620	0.172084	0.723572	1.000000	0.724320	-0.013888	0.052436	0.027492	0.038061
Smokes (packs/year)	0.131861	0.174968	-0.057013	0.092214	0.493843	0.724320	1.000000	0.001713	0.043803	0.008226	0.016292
Hormonal Contraceptives	0.029201	0.004027	-0.009563	0.116944	0.004036	-0.013888	0.001713	1.000000	0.370696	0.000188	-0.056548
Hormonal Contraceptives (years)	0.298892	0.021525	0.031976	0.221456	0.040917	0.052436	0.043803	0.370696	1.000000	0.115456	0.017955
IUD	0.279429	0.032460	-0.008826	0.198134	-0.055115	0.027492	0.008226	0.000188	0.115456	1.000000	0.749288
IUD (years)	0.215427	0.006252	-0.015697	0.148692	-0.035798	0.038061	0.016292	-0.056548	0.017955	0.749288	1.000000

	Age	Number of sexual partners	First sexual intercourse	Num of pregnancies	Smokes	Smokes (years)	Smokes (packs/year)	Hormonal Contraceptives	Hormonal Contraceptives (years)	IUD	IUD (years)
STDs	0.025241	0.055370	-0.000494	0.055698	0.111289	0.089300	0.029252	-0.045460	0.006918	0.059427	0.022691
STDs (number)	-0.001330	0.041459	0.017948	0.012938	0.100117	0.088605	0.030247	-0.053642	0.002236	0.060591	0.021286
STDs:condylomatosi	-0.013751	0.036925	0.035761	-0.031189	0.055674	0.043504	0.007599	-0.025116	0.016465	0.084794	0.034293
STDs:cervical condylomatosis	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
STDs:vaginal condylomatosis	0.009505	-0.042120	0.073945	0.002754	0.069651	0.114655	0.041939	-0.064390	-0.032782	0.035484	-0.007345
STDs:vulvo-perineal condylomatosis	-0.011499	0.038992	0.039933	-0.030813	0.058468	0.045561	0.008761	-0.029001	0.018090	0.069399	0.033032
STDs:syphilis	0.017457	0.028646	-0.094885	0.150551	0.079358	0.013850	-0.003754	-0.006252	0.003385	-0.020393	-0.026924
STDs:pelvic inflammatory disease	0.024854	0.030929	0.000256	-0.052239	-0.013974	-0.010111	-0.006901	0.023085	-0.011613	-0.011179	-0.008376
STDs:genital herpes	-0.027433	-0.031413	0.024691	-0.028411	-0.013974	-0.010111	-0.006901	0.023085	-0.016362	-0.011179	-0.008376
STDs:mollusum contagiosum	0.000722	0.030929	-0.011961	0.043074	-0.013974	-0.010111	-0.006901	-0.050547	-0.018737	-0.011179	-0.008376
STDs:AIDS	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
STDs:HIV	0.005009	0.018752	-0.007627	0.014401	0.056151	0.088930	0.053995	-0.076371	-0.035063	0.007118	0.017928
STDs:Hepatitis B	-0.027433	-0.010633	0.012473	-0.028411	0.083503	0.099313	0.101342	-0.050547	-0.018737	-0.011179	-0.008376

	Age	Number of sexual partners	First sexual intercourse	Num of pregnancies	Smokes	Smokes (years)	Smokes (packs/year)	Hormonal Contraceptives	Hormonal Contraceptives (years)	IUD	IUD (years)
<b>STDs:HPV</b>	0.040861	0.014360	0.034938	-0.023343	0.049193	0.051201	-0.008015	0.032666	0.054142	-0.015819	-0.011853
<b>STDs: Number of diagnosis</b>	-0.001606	0.053056	-0.011617	0.039195	0.090725	0.078303	0.029912	-0.062199	-0.025662	0.035791	0.012191
<b>Dx:Cancer</b>	0.110340	0.023699	0.067996	0.042765	-0.013470	0.052859	0.107229	0.011278	0.064993	0.117166	0.103148
<b>Dx:CIN</b>	0.061443	0.016669	-0.031960	-0.037752	-0.042119	-0.030476	-0.020800	-0.004397	0.003972	0.043708	0.008887
<b>Dx:HPV</b>	0.101722	0.028646	0.044727	0.054111	0.009737	0.055398	0.109118	0.028808	0.066509	0.062142	0.035869
<b>Dx</b>	0.092635	0.024597	0.036664	-0.003034	-0.069396	-0.050213	-0.034270	-0.007245	-0.008054	0.135778	0.100340
<b>Hinselmann</b>	-0.003967	-0.039098	-0.015311	0.033987	0.033333	0.070352	0.026086	0.012360	0.054264	0.052108	0.014132
<b>Schiller</b>	0.103283	-0.007230	0.005275	0.077526	0.052028	0.093479	0.017200	-0.034002	0.101250	0.096089	0.087032
<b>Citology</b>	-0.016862	0.024067	-0.009594	-0.020131	-0.004639	-0.007275	0.004250	-0.025116	0.084429	0.013292	0.007103
<b>Biopsy</b>	0.055956	-0.000408	0.008771	0.027959	0.028724	0.061204	0.024487	-0.018015	0.097937	0.059231	0.038176

In [121]:

```
# drawing a heatmap to show the correlation
plt.figure(figsize=(25, 25))
sns.heatmap(corr_matrix, annot=True)
```

Out[121]:

<AxesSubplot:>

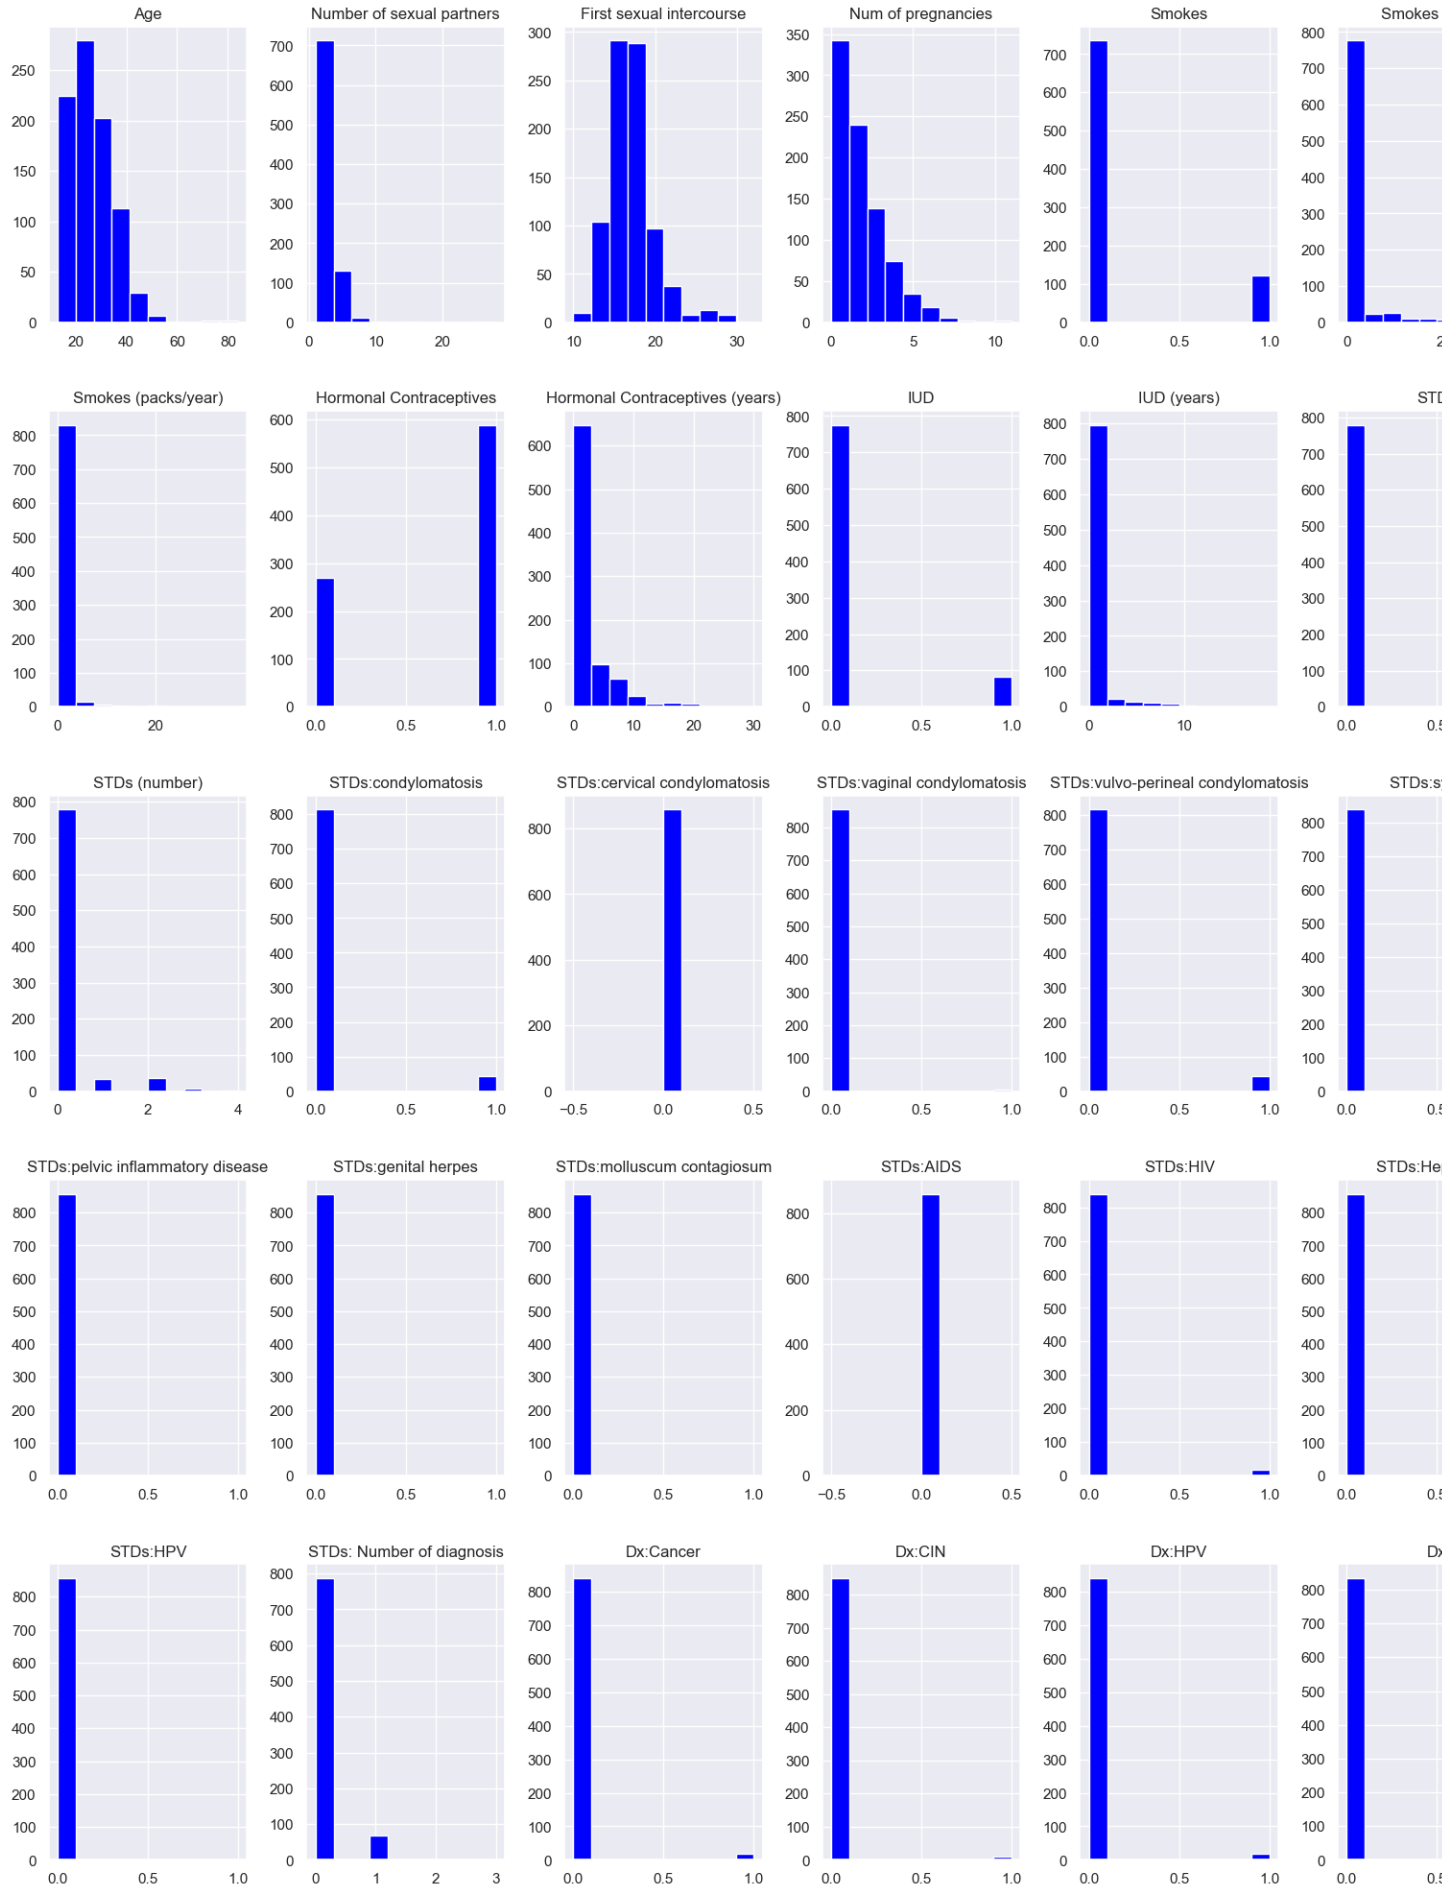


Age	1	0.086	0.37	0.51	0.057	0.22	0.13	0.029	0.3	0.28	0.2
Number of sexual partners	0.086	1	-0.15	0.082	0.24	0.18	0.17	0.004	0.022	0.032	0.001
First sexual intercourse	0.37	-0.15	1	-0.046	-0.12	-0.059	-0.057	0.009	0.032	-0.008	0.001
Num of pregnancies	0.51	0.082	-0.046	1	0.077	0.17	0.092	0.12	0.22	0.2	0.1
Smokes	0.057	0.24	-0.12	0.077	1	0.72	0.49	0.004	0.041	-0.055	-0.001
Smokes (years)	0.22	0.18	-0.059	0.17	0.72	1	0.72	-0.014	0.052	0.027	0.001
Smokes (packs/year)	0.13	0.17	-0.057	0.092	0.49	0.72	1	0.001	0.044	0.008	0.001
Hormonal Contraceptives	0.029	0.004	-0.009	0.12	0.004	-0.014	0.001	1	0.37	0.000	0.001
Hormonal Contraceptives (years)	0.3	0.022	0.032	0.22	0.041	0.052	0.044	0.37	1	0.12	0.001
IUD	0.28	0.032	-0.008	0.2	-0.055	0.027	0.008	0.000	0.19	1	0.7
IUD (years)	0.22	0.006	0.016	0.15	-0.036	0.038	0.016	-0.057	0.018	0.75	1
STDs	0.025	0.055	-0.000	0.056	0.11	0.089	0.029	-0.045	0.006	0.059	0.001
STDs (number)	0.001	0.041	0.018	0.013	0.1	0.089	0.03	-0.054	0.002	0.061	0.001
STDs:condylomatosis	-0.014	0.037	0.036	-0.031	0.056	0.044	0.007	-0.025	0.016	0.085	0.001
STDs:cervical condylomatosis											
STDs:vaginal condylomatosis	0.009	0.042	0.074	0.002	0.07	0.11	0.042	-0.064	-0.033	0.035	0.001
STDs:vulvo-perineal condylomatosis	-0.011	0.039	0.04	-0.031	0.058	0.046	0.008	-0.029	0.018	0.069	0.001

In [131]:

```
# using a histogram to visualize the whole dataset

df.hist(bins = 10, figsize = (20,30), color='blue', sharex=False)
plt.show()
```



4) PREPARING DATA BEFORE TRAINING

In [132]:  
X

Out[132]:

	Age	Number of sexual partners	First sexual intercourse	Num of pregnancies	Smokes	Smokes (years)	Smokes (packs/year)	Hormonal Contraceptives	Hormonal Contraceptives (years)	IUD	IUD (years)	STDs	Smokes (years)
0	18	4.0	15.0	1.0	0.0	0.0	0.0	0.0	0.00	0.0	0.0	0.0	0.0
1	15	1.0	14.0	1.0	0.0	0.0	0.0	0.0	0.00	0.0	0.0	0.0	0.0
2	34	1.0	15.0	1.0	0.0	0.0	0.0	0.0	0.00	0.0	0.0	0.0	0.0
3	52	5.0	16.0	4.0	1.0	37.0	37.0	1.0	3.00	0.0	0.0	0.0	0.0
4	46	3.0	21.0	4.0	0.0	0.0	0.0	1.0	15.00	0.0	0.0	0.0	0.0
...	...	...	...	...	...	...	...	...	...	...	...	...	...
853	34	3.0	18.0	0.0	0.0	0.0	0.0	0.0	0.00	0.0	0.0	0.0	0.0
854	32	2.0	19.0	1.0	0.0	0.0	0.0	1.0	8.00	0.0	0.0	0.0	0.0
855	25	2.0	17.0	0.0	0.0	0.0	0.0	1.0	0.08	0.0	0.0	0.0	0.0
856	33	2.0	24.0	2.0	0.0	0.0	0.0	1.0	0.08	0.0	0.0	0.0	0.0
857	29	2.0	20.0	1.0	0.0	0.0	0.0	1.0	0.50	0.0	0.0	0.0	0.0

858 rows × 33 columns

In [133]:  
Y  
Out[133]:

0 0  
1 0  
2 0  
3 0  
4 0  
..  
853 0  
854 0  
855 0  
856 0  
857 0

Name: Biopsy, Length: 858, dtype: int64

# scaling the data  
In [135]:

```
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
```

In [136]:

```
x_scaled = scaler.fit_transform(X)
x_scaled
```

Out[136]:

```
array([[ -1.03856336,  0.90542946, -0.70779867, ..., -0.20622158,
        -0.3072259 , -0.23249528],
       [-1.39179566, -0.91961003, -1.06544958, ..., -0.20622158,
        -0.3072259 , -0.23249528],
       [ 0.84534227, -0.91961003, -0.70779867, ..., -0.20622158,
        -0.3072259 , -0.23249528],
       ...,
       [-0.21435465, -0.31126353,  0.00750317, ..., -0.20622158,
        -0.3072259 ,  4.30116263],
       [ 0.72759817, -0.31126353,  2.51105958, ..., -0.20622158,
        -0.3072259 , -0.23249528],
       [ 0.25662176, -0.31126353,  1.08045591, ..., -0.20622158,
        -0.3072259 , -0.23249528]])
```

In [138]:

```
# splitting the data into training and testing dataset
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x_scaled, y,
test_size=0.30, random_state=355)
```

## 5) TRAINING AND EVALUATING XGBOOST CLASSIFIER

In [141]:

```
import xgboost as xgb
from xgboost import XGBClassifier
```

In [143]:

```
xgboost_model = XGBClassifier(objective='binary:logistic')
```

In [144]:

```
xgboost_model.fit(x_train, y_train)
```

Out[144]:

```
XGBClassifier(base_score=None, booster=None, callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytree=None, early_stopping_rounds=None,
              enable_categorical=False, eval_metric=None, feature_types=None,
              gamma=None, gpu_id=None, grow_policy=None,
              importance_type=None,
              interaction_constraints=None, learning_rate=None, max_bin=None,
              max_cat_threshold=None, max_cat_to_onehot=None,
              max_delta_step=None, max_depth=None, max_leaves=None,
              min_child_weight=None, missing=nan, monotone_constraints=None,
              n_estimators=100, n_jobs=None, num_parallel_tree=None,
```

```
predictor=None, random_state=None, ...)
```

In [146]:

```
# model evaluation on the training data
xgboost_model.score(x_train, y_train)
```

Out[146]:

```
0.9983333333333333
```

In [158]:

```
# model evaluation on the testing data
xgboost_model.score(x_test, y_test)
```

Out[158]:

```
0.9496124031007752
```

the valiation between the training data and the test data, seems like my model has overfitted on the training data

## hyperparameter tuning

In [163]:

```
from sklearn.model_selection import GridSearchCV
param_grid = {
    'learning_rate':[0.1],
    'max_depth' : range(20),
    'n_estimators' : range(100)
}
```

In [164]:

```
grid_search = GridSearchCV(XGBClassifier(objective='binary:logistic'),
param_grid, verbose=3)
grid_search.fit(x_train, y_train)
Fitting 5 folds for each of 2000 candidates, totalling 10000 fits
[CV 1/5] END learning_rate=0.1, max_depth=0, n_estimators=0;, score=0.942
total time= 0.0s
[CV 2/5] END learning_rate=0.1, max_depth=0, n_estimators=0;, score=0.942
total time= 0.0s
[CV 3/5] END learning_rate=0.1, max_depth=0, n_estimators=0;, score=0.933
total time= 0.0s
[CV 4/5] END learning_rate=0.1, max_depth=0, n_estimators=0;, score=0.933
total time= 0.0s
[CV 5/5] END learning_rate=0.1, max_depth=0, n_estimators=0;, score=0.933
total time= 0.0s
[CV 1/5] END learning_rate=0.1, max_depth=0, n_estimators=1;, score=nan total
time= 0.0s
[CV 2/5] END learning_rate=0.1, max_depth=0, n_estimators=1;, score=nan total
time= 0.0s
[CV 3/5] END learning_rate=0.1, max_depth=0, n_estimators=1;, score=nan total
time= 0.0s
[CV 4/5] END learning_rate=0.1, max_depth=0, n_estimators=1;, score=nan total
time= 0.0s
```

Out[164]:

```
GridSearchCV(estimator=XGBClassifier(base_score=None, booster=None,
                                     callbacks=None, colsample_bylevel=None,
                                     colsample_bynode=None,
                                     colsample_bytree=None,
                                     early_stopping_rounds=None,
                                     enable_categorical=False,
                                     eval_metric=None,
                                     feature_types=None, gamma=None,
                                     gpu_id=None, grow_policy=None,
                                     importance_type=None,
                                     interaction_constraints=None,
                                     learning_rate=None, max_bin=None,
                                     max_cat_threshold=None,
                                     max_cat_to_onehot=None,
                                     max_delta_step=None, max_depth=None,
                                     max_leaves=None, min_child_weight=None,
                                     missing=nan, monotone_constraints=None,
                                     n_estimators=100, n_jobs=None,
                                     num_parallel_tree=None, predictor=None,
                                     random_state=None, ...),
             param_grid={'learning_rate': [0.1], 'max_depth': range(0, 20),
                         'n_estimators': range(0, 100)},
             verbose=3)
```

In [167]:

```
grid_search.best_params_
```

Out[167]:

```
{'learning_rate': 0.1, 'max_depth': 1, 'n_estimators': 63}
```

In [168]:

```
new_xgboost_model = XGBClassifier(learning_rate=0.1, max_depth=1,
                                   n_estimators=63)
new_xgboost_model.fit(x_train, y_train)
```

Out[168]:

```
XGBClassifier(base_score=None, booster=None, callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytree=None, early_stopping_rounds=None,
              enable_categorical=False, eval_metric=None, feature_types=None,
              gamma=None, gpu_id=None, grow_policy=None,
              importance_type=None,
              interaction_constraints=None, learning_rate=0.1, max_bin=None,
              max_cat_threshold=None, max_cat_to_onehot=None,
              max_delta_step=None, max_depth=1, max_leaves=None,
              min_child_weight=None, missing=nan, monotone_constraints=None,
              n_estimators=63, n_jobs=None, num_parallel_tree=None,
              predictor=None, random_state=None, ...)
```

```
# evaluation of the new model on the training data
new_xgboost_model.score(x_train, y_train)
```

In [170]:

```
0.965
```

Out[170]:

```
# evaluation of the new model on the test data
new_xgboost_model.score(x_test, y_test)
```

In [171]:

```
0.9651162790697675
```

Out[171]:

my test and the train score is almost same, now i can say that my model is performing better and has not overfitted on the training data

```
# evaluating the model using a confusion matrix
from sklearn.metrics import confusion_matrix, classification_report
y_pred = new_xgboost_model.predict(x_test)
cof_mat = confusion_matrix(y_test, y_pred)
cof_mat
```


In [186]:

```
array([[235,    6],
       [  3,   14]], dtype=int64)
```

Out[186]:

```
# drawing a confusion matrix using heat map
sns.heatmap(cof_mat, annot=True)
plt.show()
```

In [187]:





In [188]:

```
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.99	0.98	0.98	241
1	0.70	0.82	0.76	17
accuracy			0.97	258
macro avg	0.84	0.90	0.87	258
weighted avg	0.97	0.97	0.97	258

## 6) saving new\_xgb\_model

In [189]:

```
import pickle
with open('new_xgboost_model', 'wb') as file:
    pickle.dump(new_xgboost_model, file)
    file.close
```

In [190]:

```
# saving the standardscaler model
with open('StandardScaler.pickle', 'wb') as file:
    pickle.dump(scaler, file)
    file.close()
```

In [ ]: